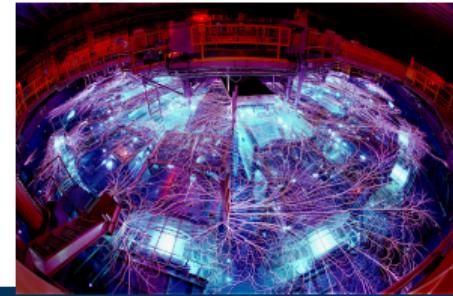


Exceptional service in the national interest



Trilinos Tutorial

Christian Glusa and Chris Siefert {caglusa,csiefer}@sandia.gov

Presented at ICPP
September 11, 2025

Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525. SAND NO. SAND2025-11032C

TRILINOS AREAS



DevSecOps:

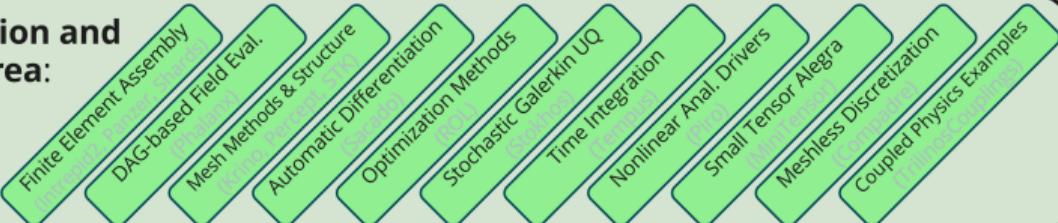


Sam Browne



Mauro Perego

Discretization and Analysis Area:



Jonathan Hu

Solvers Area:

Solver Interface (Stratimikos)

Nonlinear Solvers (NOX)

Krylov Solvers (Belos)

Eigen Solvers
(Anasazi)

Direct Sparse &
Dense Solvers
(Amesos2, Adelus)

Block Solvers
(Teuchos)

Multigrid Methods
(MueLu)

Multilevel Domain Decomposition Methods (ShyLU)

Iterative Substructuring
(ShyLU)

Overlapping Schwarz
(ShyLU/Trilinos)

Polynomial Smoothers, SOR, ILU (ifpack2)

Linear Factorization
(KLU, ShyLU/Basket)

Cholesky and LDL Factorizations
(ShyLU/Tacho)



Roger Pawlowski

Core Area:

Python Interface (PyTrilinos2)

Abstract Numerical Algorithms & Interfaces (Thyra)

Matrix Generation
Tools (Galeri)

Vector Reduction/Transformation
Operators (RTOp)

Distributed Linear Algebra
(Tpetra)

Common SE Utilities (Teuchos)

Optimized Math Kernels (Kokkos Kernels)

Performance Portability (Kokkos)

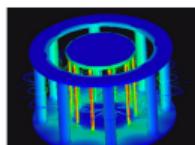
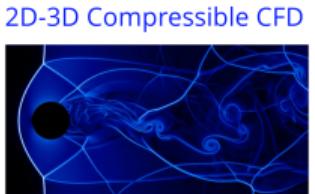
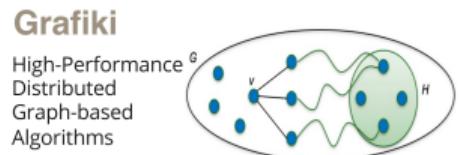
Finite Element
Utilities (SEACAS)

Load balancing
(Zoltan/Zoltan2)

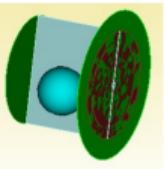
Mesh Generation
(Parmetis)

COLLABORATIONS WITH USERS & STAKEHOLDERS

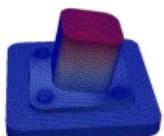
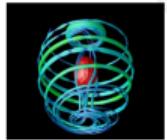
Trilinos currently has collaborations with ~25 applications ...



Gemma
High-fidelity, robust, solutions for Maxwell's Equations



deal.II
Rapid development of modern finite element codes



Instructions

■ Connect to AWS

- Go to bit.ly/kokkos-vms and claim an IP
- Download the linked ssh key and set its permissions:
`chmod 400 kokkos-tutorials.pem`
- Log into them with ssh:

```
ssh -i kokkos-tutorial.pem ec2-user@<YOUR_IP> -o PubkeyAcceptedAlgorithms=+ssh-rsa
```

■ `git clone https://github.com/trilinos/icpp.git`

■ Launch the Trilinos container:

```
cd icpp  
./containerLauncher.sh
```

■ In the container:

- Trilinos build directory:
`/workspace/trilinos/build/`
- Location of Tpetra examples
`/workspace/trilinos/build/packages/tpetra/core/example`
- Location of Belos/Lfpack2/MueLu example
`/workspace/trilinos/build/packages/muelu/example/basic`

Distributed Sparse Linear Algebra

Tpetra Tutorial

Go to <https://docs.trilinos.org/dev/packages/tpetra/doc/html/index.html>

Iterative Solvers

Discretization of partial differential equations gives rise to large linear systems of equations

$$\mathbf{A}\vec{x} = \vec{b},$$

where \mathbf{A} is sparse, i.e. only a few non-zero entries per row.

Example

2D Poisson equation:

$$\begin{aligned} -\Delta u &= f \text{ in } \Omega = [0, 1]^2, \\ u &= 0 \text{ on } \partial\Omega. \end{aligned}$$

Central finite differences on a uniform mesh $\{x_{i,j}\}$:

$$\begin{aligned} 4u_{i,j} - u_{i,j+1} - u_{i,j-1} - u_{i+1,j} - u_{i-1,j} &= f(x_{i,j})\Delta x^2 && \text{if } x_{i,j} \notin \partial\Omega, \\ u_{i,j} &= 0 && \text{if } x_{i,j} \in \partial\Omega. \end{aligned}$$

→ 5 entries or less per row of \mathbf{A} .

Instead of dense format, keep matrix \mathbf{A} in a sparse format e.g. *compressed sparse row* (CSR):

$$\mathbf{A} = \begin{pmatrix} 1 & 2 & 0 \\ 3 & 4 & 0 \\ 0 & 0 & 5 \end{pmatrix}$$

$$\text{rowptr} = (\textcolor{red}{0} \textcolor{green}{2} \textcolor{blue}{4} \textcolor{cyan}{5})$$

$$\text{indices} = (\textcolor{red}{0} \textcolor{green}{1} \textcolor{blue}{0} \textcolor{cyan}{1} \textcolor{magenta}{2})$$

$$\text{values} = (\textcolor{red}{1} \textcolor{green}{2} \textcolor{blue}{3} \textcolor{cyan}{4} \textcolor{magenta}{5})$$

Available solvers

Solve

$$\mathbf{A}\vec{x} = \vec{b}.$$

Option 1: Direct solvers (think Gaussian elimination)

- Factorisation scales as $\mathcal{O}(n^3)$.
- Factors are a lot denser than \mathbf{A} → memory cost.
- Parallel implementation not straightforward.
- Does not require a lot of information about the structure of \mathbf{A} .

Observation

\mathbf{A} has $\mathcal{O}(n)$ non-zero entries. → Optimal complexity for a solve is $\mathcal{O}(n)$ operations.

Option 2: Iterative solvers

- Exploit an operation that has $\mathcal{O}(n)$ complexity: mat-vec.
- Easy to parallelize.
- Can have small memory footprint. (In the best case, we only need to keep a single vector.)
- Generally more restrictions on properties of \mathbf{A} .

Available solvers

Solve

$$\mathbf{A}\vec{x} = \vec{b}.$$

Option 1: Direct solvers (think Gaussian elimination)

- Factorisation scales as $\mathcal{O}(n^3)$.
- Factors are a lot denser than \mathbf{A} → memory cost.
- Parallel implementation not straightforward.
- Does not require a lot of information about the structure of \mathbf{A} .

Observation

\mathbf{A} has $\mathcal{O}(n)$ non-zero entries. → Optimal complexity for a solve is $\mathcal{O}(n)$ operations.

Option 2: Iterative solvers

- Exploit an operation that has $\mathcal{O}(n)$ complexity: mat-vec.
- Easy to parallelize.
- Can have small memory footprint. (In the best case, we only need to keep a single vector.)
- Generally more restrictions on properties of \mathbf{A} .

Krylov methods

Based on mat-vecs, we can compute

$$\vec{y}^0 := \vec{b}$$

("initial guess")

$$\vec{y}^{k+1} := \vec{y}^k + \underbrace{\left(\vec{b} - \mathbf{A}\vec{y}^k \right)}_{\text{"residual"}}$$

and recombine in some smart way to obtain an approximate solution

$$\bar{x}^K = \sum_{k=0}^K \alpha_k \vec{y}^k.$$

Expressions for α_k typically involve inner products between vectors in the so-called *Krylov space*
 $\text{span } \{\vec{y}^k\} = \left\{ \vec{b}, \mathbf{A}\vec{b}, \mathbf{A}^2\vec{b}, \mathbf{A}^3\vec{b}, \dots, \mathbf{A}^K\vec{b} \right\}.$

- Keeping the entire Krylov space can be quite expensive.
- Computing inner products involves an all-reduce which can be costly at large scale.

Two particular Krylov methods:

- Conjugate gradient (CG)
 - Use a short recurrence, i.e. does not keep the whole Krylov space around.
 - Provably works for symmetric positive definite (spd) \mathbf{A} .
- Generalized Minimum Residual (GMRES, GMRES(K))
 - Works for nonsymmetric systems.
 - GMRES keeps the whole Krylov space around.
 - GMRES(K) discards the Krylov space after K iterations.

Convergence of Krylov methods

CG convergence result:

$$\|\vec{x}^K - \vec{x}\| \leq \left(1 - 1/\sqrt{\kappa(\mathbf{A})}\right)^K \|\vec{x}^0 - \vec{x}\|,$$

where $\kappa(\mathbf{A})$ is the *condition number* of \mathbf{A} :

$$\kappa(\mathbf{A}) = \|\mathbf{A}\| \|\mathbf{A}^{-1}\|.$$

A common theme with Krylov methods:

κ measures how hard it is to solve the system, i.e. how many iterations are required to reach a given tolerance.

Idea

Reduce the condition number ("Preconditioning").

Instead of solving

$$\mathbf{A}\vec{x} = \vec{b},$$

solve

$$\mathbf{P}\mathbf{A}\vec{x} = \mathbf{P}\vec{b}$$

or

$$\mathbf{A}\mathbf{P}\vec{z} = \vec{b}, \quad \vec{x} = \mathbf{P}\vec{z}$$

with *preconditioner* \mathbf{P} so that $\kappa(\mathbf{P}\mathbf{A}) \ll \kappa(\mathbf{A})$.

Two requirements that must be balanced:

- Multiplication with \mathbf{P} should be comparable in cost to \mathbf{A} .
- $\mathbf{P} \approx \mathbf{A}^{-1}$.

Some simple preconditioners

- Jacobi: $\mathbf{P} = \mathbf{D}^{-1}$, where \mathbf{D} is the diagonal of \mathbf{A} .
- Gauss-Seidel: $\mathbf{P} = (\mathbf{D} + \mathbf{L})^{-1}$, where \mathbf{L} is the lower or upper triangular part of \mathbf{A} .
- Polynomial preconditioners: $\mathbf{P} = p(\mathbf{A})$, where p is some carefully chosen polynomial.
- Incomplete factorizations such as ILU or Incomplete Cholesky.

Krylov methods and preconditioners: Packages in the Trilinos project



- Support for hybrid (MPI+X) parallelism,
 $X \in \{\text{OpenMP, CUDA, HIP, ...}\}$
- C++, open source

Belos - iterative linear solvers

- Standard methods:
 - Conjugate Gradients (CG), Generalized Minimal Residual (GMRES)
 - TFQMR, BiCGStab, MINRES, Richardson / fixed-point
- Advanced methods:
 - Block GMRES, block CG/BiCG
 - Hybrid GMRES, GCRODR (block recycling GMRES)
 - TSQR (tall skinny QR), LSQR
- Ongoing research:
 - Communication avoiding methods
 - Pipelined and s-step methods
 - Mixed precision methods

Ifpack2 - single-level solvers and preconditioners

- incomplete factorisations
 - ILUT
 - RILU(k)
- relaxation preconditioners
 - Jacobi
 - Gauss-Seidel (and a multithreaded variant)
 - Successive Over-Relaxation (SOR)
 - Symmetric versions of Gauss-Seidel and SOR
 - Chebyshev
- additive Schwarz domain decomposition

Hands-on: Krylov methods and preconditioning

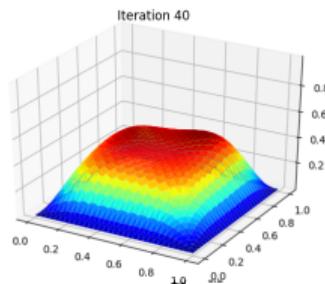
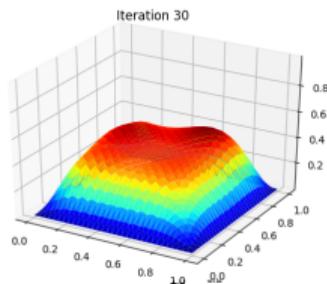
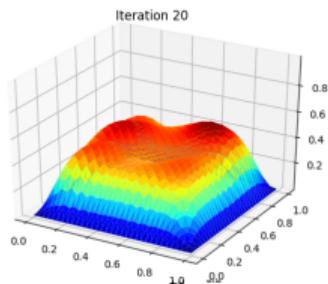
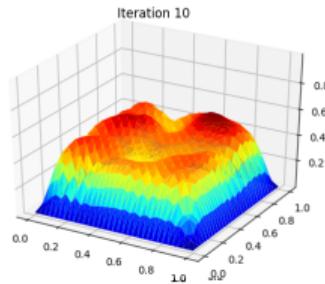
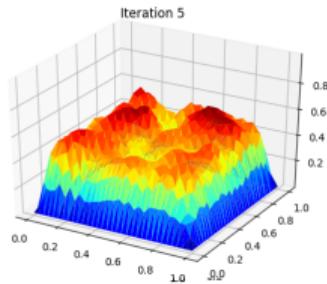
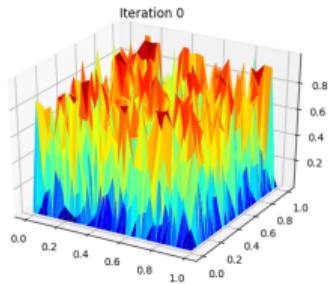
Go to <https://github.com/trilinos/icpp/blob/main/lesson/lesson.md>

Sets 1 and 2

Motivation for Multigrid methods

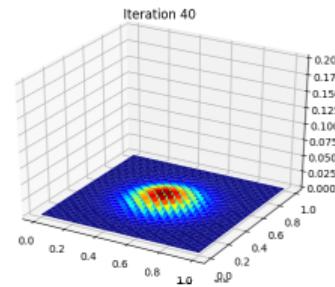
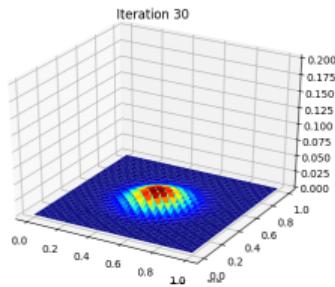
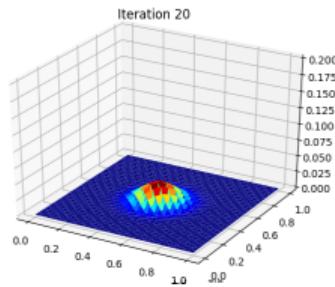
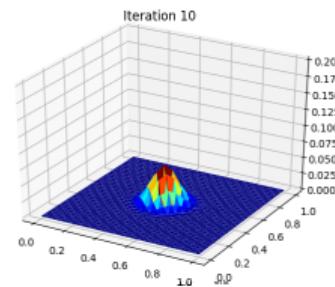
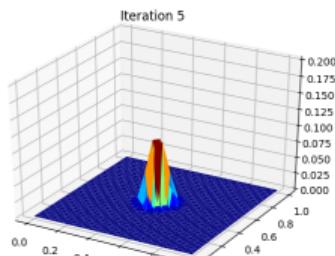
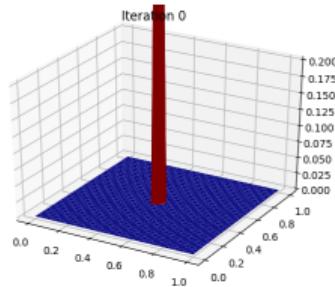
Convergence of Jacobi: $\vec{y}^{k+1} = \vec{y}^k + \mathbf{D}^{-1}\vec{r}^k$, $\vec{r}^k = \vec{b} - \mathbf{A}\vec{y}^k$

High frequency error is damped quickly, low frequency error slowly



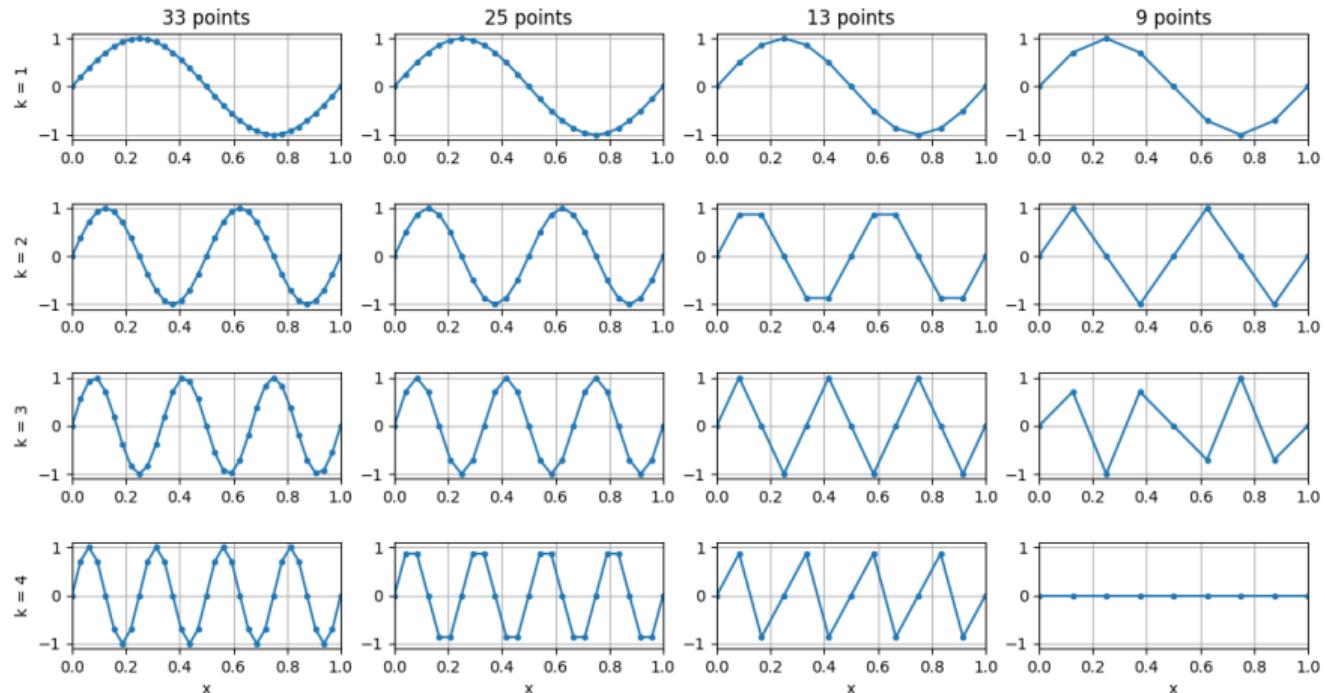
Motivation for Multigrid methods

Convergence of Jacobi:
Local transmission of information cannot result in a scalable method



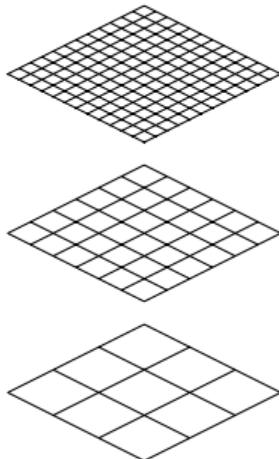
Motivation for Multigrid methods

Resolution affects observed frequency:



Idea: accelerate Jacobi convergence by reducing resolution!

Multigrid



- Main idea: accelerate solution of $\mathbf{A}\vec{x} = \vec{b}$ by using "hierarchy" of coarser problems
- Remove high-frequency error on fine mesh, where application matrix lives (using Jacobi or another cheap preconditioner),
- Move to coarser mesh
- Remove high-frequency error on coarser mesh by solving residual equation
- Move to coarser mesh
- :
- Solve a small problem on a very coarse mesh.
- Move back up.

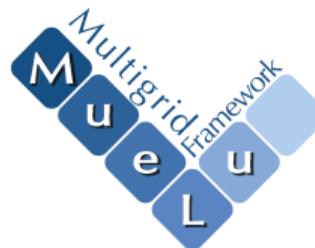
Repeat.

- *Geometric multigrid* requires coarse mesh information.
- *Algebraic multigrid* constructs coarser matrices on the fly based on fine-level matrix entries.

The Trilinos/MueLu package

- Algebraic Multigrid package in Trilinos

Templated C++ library with support for 2B+ unknowns and next-generation architectures (OpenMP, CUDA, HIP, ...)
- Robust, scalable, portable AMG preconditioning is critical for many large-scale simulations
 - Multifluid plasma simulations
 - Shock physics
 - Magneto-hydrodynamics (MHD)
 - Low Mach computational fluid dynamics (CFD)
- Capabilities
 - Aggregation-based and structured coarsening
 - Smoothers: Jacobi, Gauss-Seidel, ℓ_1 Gauss-Seidel, multithreaded Gauss-Seidel, polynomial, ILU
 - Load balancing for good parallel performance
- Ongoing research
 - performance on next-generation architectures
 - AMG for multiphysics
 - Multigrid for coupled structured/unstructured problems
 - Algorithm selection via machine learning

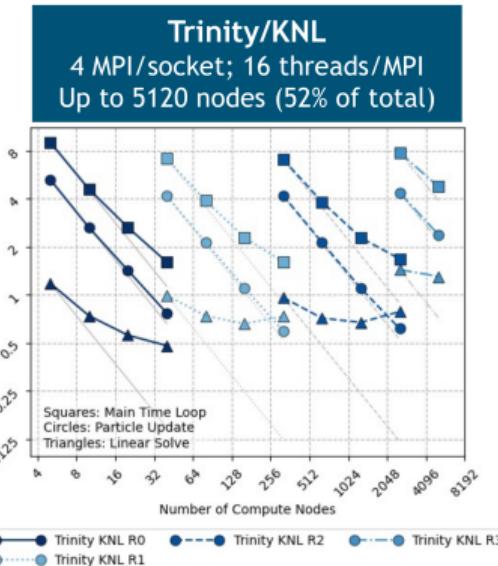
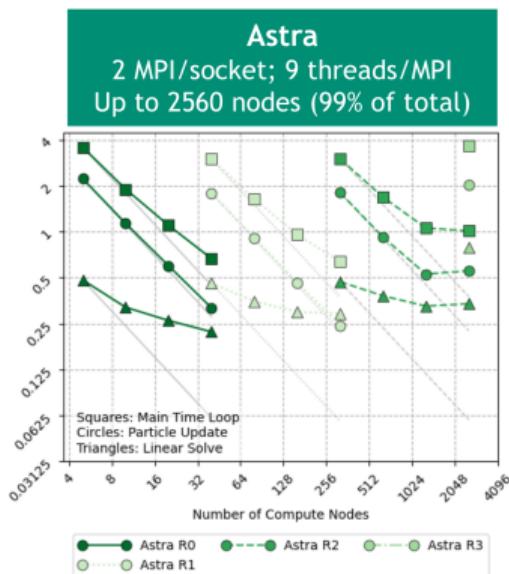
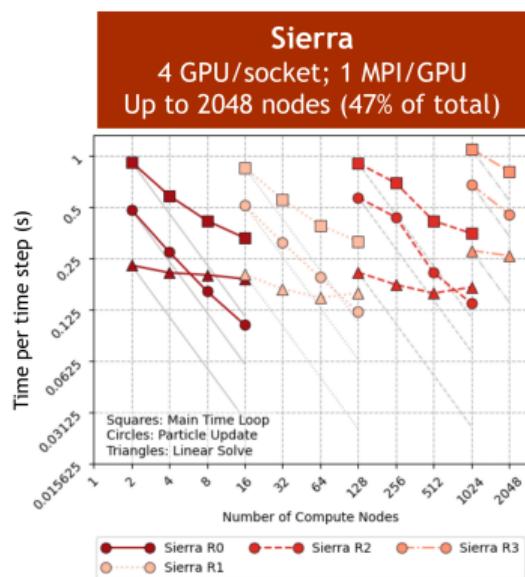


Hands-on: Algebraic Multigrid

Go to <https://github.com/trilinos/icpp/blob/main/lesson/lesson.md>
Set 3 & 4

Strong & weak scaling results for EMPIRE (Maxwell + PIC)

- Specialized multigrid for curl-curl problem
- Largest problem to date: 34B unknowns



Mesh	Elements	Nodes	Edges	Particles
R0	3.7M	660k	4.4M	360M
R1	25M	4.4M	30M	2.4B
R2	200M	32M	240M	19B
R3	1.6B	270M	1.9B	160B

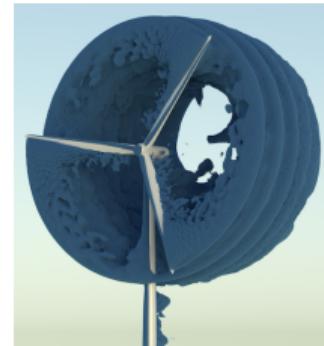
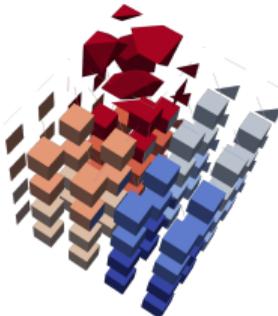
Ongoing work

- Multiprecision (Krylov methods with mixed precision; lower precision preconditioning)
- Multigrid approaches for higher order discretizations
- Matrix-free multigrid
- Multigrid on semi-structured meshes
- Machine learning for AMG coarsening
- Preconditioning for multiphysics systems
- Multigrid for hierarchical matrices (boundary integral and nonlocal equations)

Algorithm 1 Iterative Refinement with GMRES Error Correction

```

1:  $r_0 = b - Ax_0$  [double]
2: for  $i = 1, 2, \dots$  until convergence: do
3:   Use GMRES( $m$ ) to solve  $Au_i = r_i$  for correction  $u_i$  [single]
4:    $x_{i+1} = x_i + u_i$  [double]
5:    $r_{i+1} = b - Ax_{i+1}$  [double]
6: end for
```



Take away messages

- CG works for spd matrix and preconditioner.
- GMRES works for unsymmetric systems, but requires more memory.
- Simple preconditioners can reduce the number of iterations, but often do not lead to a scalable solver.
- Multigrid (when applicable) has constant number of iterations, independent of the problem size.

Thank you for your attention!

Interested in working on Multigrid (and other topics) at a national lab?

We are always looking for motivated

- summer and year-round students,
- postdocs.

Please contact us!